

For more details, please visit <https://assistedlabs.com>

Cost Optimization in OpenShift – Quotas, Limits, and Chargeback Models

Cost Optimization in OpenShift – Quotas, Limits, and Chargeback Models



Introduction

OpenShift, Red Hat's Kubernetes-based enterprise container platform, enables organizations to build, deploy, and manage applications at scale. While it provides agility, flexibility, and scalability, these benefits often come with a challenge: **managing and optimizing costs**.

In shared or multi-tenant clusters, resource usage by different teams can quickly escalate if not governed properly. Cost optimization is not just about reducing expenses but also about ensuring that resources are **efficiently allocated, fairly distributed, and transparently accounted for**.

Three key mechanisms in OpenShift that contribute to cost optimization are:

To get the OpenShift Hands-on Experience, kindly visit: <https://assistedlabs.com>

ASSISTED LABS

For more details, please visit <https://assistedlabs.com>

1. **Quotas** – restricting total resource usage at the project/namespace level.
2. **Limit Ranges** – enforcing container-level minimums and maximums.
3. **Chargeback Models** – enabling financial accountability by mapping usage to costs.

This document explores these three mechanisms in depth and provides best practices for organizations to achieve sustainable cost efficiency in OpenShift environments.

Resource Quotas in OpenShift

What Are Resource Quotas?

A **Resource Quota** defines constraints that limit the aggregate resource consumption in a namespace. For example, administrators can restrict CPU, memory, storage, or even the number of objects like Pods, Services, and ConfigMaps.

Why Are Quotas Important?

- **Prevent Resource Hogging:** Avoids scenarios where one team consumes the majority of resources, leaving others starved.
- **Promotes Fair Usage:** Ensures equitable distribution across teams and projects.
- **Improves Planning:** Helps developers understand resource boundaries and design their workloads accordingly.

Examples of Quotas

- Maximum of 10 CPUs per namespace.
- Memory restricted to 32 GB.
- Maximum of 20 Pods per namespace.
- Limiting Persistent Volume Claims (PVCs) to 5 per project.

Implementation

Resource quotas are configured using YAML manifests. For example:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources
spec:
  hard:
    requests.cpu: "10"
    requests.memory: 32Gi
    limits.cpu: "20"
```

ASSISTED LABS

For more details, please visit <https://assistedlabs.com>

```
limits.memory: 64Gi
```

This ensures that all Pods combined in the namespace cannot exceed the set boundaries.

Limit Ranges in OpenShift

What Are Limit Ranges?

A **Limit Range** provides default, minimum, and maximum values for container-level resource requests and limits. Unlike quotas that apply cluster-wide within a namespace, limit ranges focus on individual workloads.

Benefits of Limit Ranges

- **Prevents Overconsumption:** Stops a single container from monopolizing cluster resources.
- **Encourages Right-Sizing:** Developers must explicitly request realistic resources.
- **Automates Defaults:** Containers without specified requests automatically inherit default values.

Examples

- Minimum memory per container: 128 MB.
- Maximum CPU per container: 2 cores.
- Default memory assigned: 512 MB if not specified.

Implementation Example

```
apiVersion: v1
kind: LimitRange
metadata:
  name: mem-limit-range
spec:
  limits:
  - default:
      memory: 512Mi
    defaultRequest:
      memory: 256Mi
    type: Container
```

This enforces boundaries while ensuring applications have predictable resource usage.

To get the OpenShift Hands-on Experience, kindly visit: <https://assistedlabs.com>

For more details, please visit <https://assistedlabs.com>

Chargeback Models in OpenShift

What Is Chargeback?

A **Chargeback Model** assigns costs to teams or business units based on their resource usage. It creates accountability, promotes responsible consumption, and helps IT teams recover costs for infrastructure usage.

Benefits

- **Transparency:** Teams understand how their consumption translates into expenses.
- **Efficiency:** Encourages developers to use resources more responsibly.
- **Budgeting & Forecasting:** Provides accurate financial data for planning future infrastructure needs.

Chargeback Approaches

1. **Fixed Allocation** – Each team gets a pre-defined resource allocation and corresponding budget.
2. **Usage-Based Chargeback** – Teams are billed based on actual usage (CPU hours, memory GBs, storage, etc.).
3. **Hybrid Model** – Combines baseline allocations with variable costs for extra consumption.

Implementation in OpenShift

- Use **OpenShift Metering**, **Prometheus**, or external billing tools to track usage.
 - Map metrics (CPU, memory, storage) to cost rates.
 - Generate monthly or quarterly reports to share with stakeholders.
-

Best Practices for Cost Optimization

Organizations can significantly reduce unnecessary expenses and improve operational efficiency by combining quotas, limits, and chargeback models.

To get the OpenShift Hands-on Experience, kindly visit: <https://assistedlabs.com>

ASSISTED LABS

For more details, please visit <https://assistedlabs.com>

1. **Set Appropriate Quotas**
 - Base quotas on actual business requirements.
 - Review regularly to ensure fairness.
2. **Apply Limit Ranges**
 - Enforce minimum and maximum boundaries.
 - Prevent “noisy neighbor” issues.
3. **Implement Chargeback**
 - Establish clear financial accountability.
 - Use dashboards and reports for visibility.
4. **Monitor Continuously**
 - Tools like **Prometheus + Grafana** provide real-time insights.
 - Regularly tune allocations.
5. **Educate Developers**
 - Encourage awareness of cost implications.
 - Train teams on efficient resource usage.
6. **Review and Adjust**
 - Update policies as workloads evolve.
 - Balance performance needs with cost control.

Conclusion

Cost optimization in OpenShift is a **continuous process**, not a one-time exercise. By leveraging **Quotas** and **Limit Ranges**, administrators can enforce controlled and fair usage of cluster resources. Implementing **Chargeback Models** ensures financial accountability and promotes transparency across teams.

When combined with effective monitoring and awareness, these strategies enable organizations to run workloads **more efficiently, sustainably, and cost-effectively**.

Ultimately, cost optimization in OpenShift aligns technical operations with business objectives—delivering both operational efficiency and financial discipline.

To get the OpenShift Hands-on Experience, kindly visit: <https://assistedlabs.com>